

dm24s1

Topic 09 : Regression2

Part 01 : Overview

Dr Bernard Butler

Department of Computing and Mathematics, WIT.
(bernard.butler@setu.ie)

Autumn Semester, 2024

Outline

- Regression as a means of minimising sum of the squared errors
- Regression assumptions - what they mean, how they can be used for validation and model building
- Case studies from Advertising and Credit Balance prediction

Data Mining (Week 9)


Introduction




Motivating Example

Preparation

Data Handling



Exploring Data 1




Exploring Data 2



Building Models

Prediction

Clustering




Regression
1



Classification
1




Regression
2



Classification
2

Wrap up



Overview — Summary

1. Introduction
2. Regression1 review
3. Case Study 1: Generated
4. Case Study 3: Advertising
5. Case Study 4: Credit Balances
6. Multivariate Analysis
7. Review and Summary
8. Resources

This Week's Aim

This week's aim is to continue the introduction to linear regression, focusing more on how to deal with problems with more challenging datasets.

- Examine some extensions to the simplest case of linear regression.
- We introduce two new concepts: dimensionality reduction and regularisation
- To provide context we will use the following datasets:
 - Generated data (various)
 - Advertising dataset: predicting widgets sold based on spending in different advertising channels
 - Credit dataset: predicting credit balance using income, status, etc.

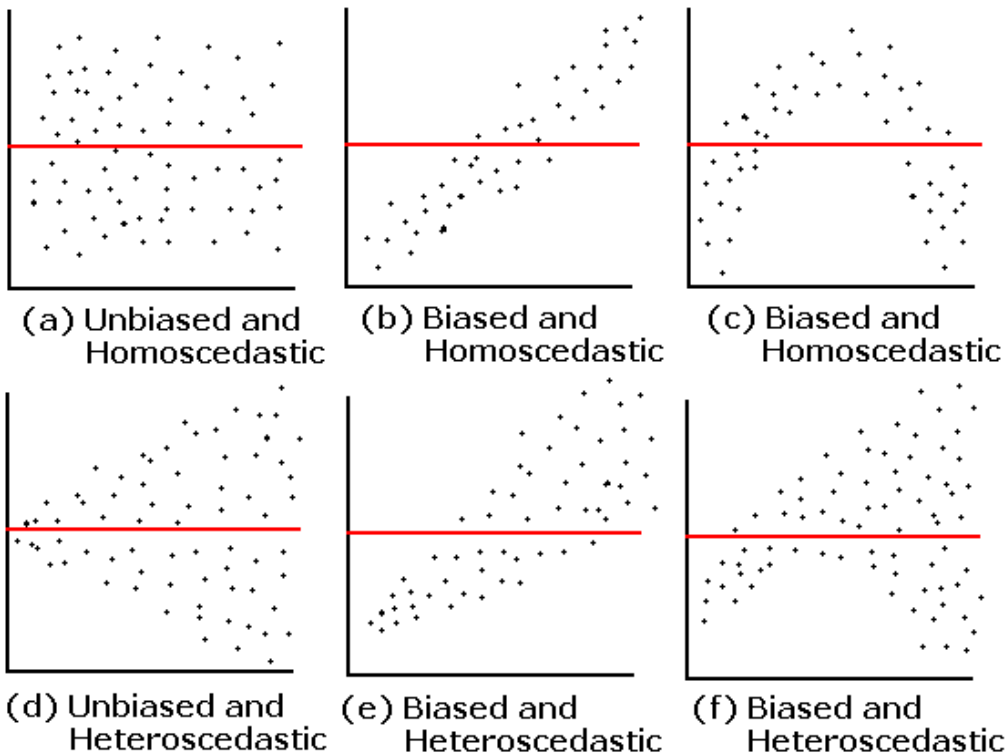
Assumptions required for the linear model to be meaningful

Definition 1 (Linear Regression Assumptions)

- ① The underlying relationship between the predictors and the response is linear in the regression parameters β .
- ② The residual errors ϵ are drawn from a (multivariate) Normal distribution $N(\mu, \sigma^2)$ where $\mu = \mathbf{0}$.
- ③ The predictors are not pairwise collinear, i.e., each pair of predictors β_{j_1} and β_{j_2} (associated with columns $X(:, j_1)$ and $X(:, j_2)$) have low correlation (equivalently, the inner product of $X(:, j_1)$ and $X(:, j_2)$ is far from zero).
- ④ There is no auto-correlation in \mathbf{y} : each observation is independent of its “neighbours”.
- ⑤ The errors are *homoscedastic* (i.e., $\text{Var}(\epsilon)$ is constant over the range of \mathbf{x} or \mathbf{y}).

➤ These assumptions can be used constructively, when model building, or as checks, when validating models.


Bias and variance in regression



Source: <https://bit.ly/3vC9zK7>


- Bias is caused by underfitting.
- Fix bias by adding suitable predictors.
- Overfitting causes large variance.
- If variance changes over the range, some errors get undue attention.
- Fix this by weighting the errors so the weighted errors satisfy $w_i e_i \approx w_j e_j, \forall i, j$.
- In practice, $w_i \approx \frac{1}{\widehat{\text{Var}}(e_i)}$.
 - Using scikit-learn: add the argument `sample_weight = someWeights`, e.g., `model.fit(Xtrain, yTrain, sample_weight=someWeights)`.
 - Using statsmodels: use the weighted version of least squares: `WLS(y, X, someWeights)` not `OLS(y, X)`

What's happening here???




- 1 Data is quite noisy
- 2 Training data has gaps near the edges
- 3 Model may be overfitting

Case Study 1: Heteroscedasticity - Step 1




I generated 41 x, y points based on $y = 3 - 2x$, but with added errors that increase away from $x = 0$. The plot shows the line with $\beta = (3, -2)$ in grey.

Case Study 1: Heteroscedasticity - Step 2




In this plot I added the OLS fit as a dashed line. Note that the parameters of the fit are quite different:
 $\beta_{OLS} \approx (2, -6)$.

Case Study 1: Heteroscedasticity - Step 3



This plot shows how the OLS residuals ϵ_{OLS} increase rapidly away from 0, as expected (since this was how the data was generated).

Case Study 1: Heteroscedasticity - Step 4



By inspecting the previous residual plot I estimated a weighting function so that the residuals would be “more constant”. When this was used to scale the residuals, the resulting Weighted Least Squares estimates were $\beta \approx (2.6, -4.5)$ (shown as a dotted line) and hence closer to the “true” β .

Can you see a problem with finding the weights?

If the weights are computed from the errors, they depend on the fit, hence on the weights!!


Case Study 3: Advertising: Data and Hypotheses

TV Radio Newspaper Sales

0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9


In this data set, the sales figure captures thousands of widgets of a particular type wss sold in a year. Newspaper, Radio and TV represent the annual spend per widget type on the associated advertising channel. The hypothesis is that spend on advertising is a good predictor of sales performance. Since marketing budgets are limited, where should the adverts be placed for maximum sales?

Case Study 3: Advertising: Looking at the data



Which of the advertising channels appear to have a linear relationship with Sales?

Case Study 3: Advertising: Collinearity?



Correlation matrix can indicate which attributes should participate in the model as predictors.

A good predictor should have a high correlation with the dependent variable (Sales in this case) and should have low correlation with other candidate predictors.

What are expected to be good predictors for this data?

Sidebar: specifying models

The statsmodels way

- The dataframe contains the observed variables
- The model is specified separately
- Easier to change the model when experimenting

The sklearn way

- The dataframe contains the (computed) features
- The model is defined implicitly
- Standard interface across all sklearn

statsmodels models look like "Sales ~ TV * Radio + poly(Newspaper,2)"; notation came from applied statistics community.

statsmodels offers its own plotting (like seaborn but not as good). Model summary is very convenient.

sklearn exposes more of the details (e.g., choice of algorithm and configuration parameters).

Both statsmodels and sklearn use the same libraries (scipy, numpy, etc.) underneath.

Case Study 3: Advertising: Model Building

- Start from a “full model” and prune, versus from an “empty model” and add
- We choose the latter, as it is often easier to avoid overfitting


Example 2 (Forward Selection for Advertising Data)

Define: model score: mean-square-error on the test set for a given model.

- 1 Fit “Sales \sim Newspaper”, “Sales \sim Radio”, “Sales \sim TV” and calculate their loss values.
- 2 Choose the best (lowest loss) single-term model (“Sales \sim TV” in this case), with loss $\text{MSE}(\text{TV})$.
- 3 Fit “Sales \sim TV + Newspaper” and “Sales \sim TV + Radio” and choose the lowest loss score, which is “Sales \sim TV + Radio” with loss being $\text{MSE}(\text{TV} + \text{Radio})$, which is significantly better.
- 4 Fit “Sales \sim TV + Radio + Newspaper”. Its loss is the same ($\text{MSE}(\text{TV} + \text{Radio}) \approx \text{MSE}(\text{TV} + \text{Radio} + \text{Newspaper})$), so we favour the existing simpler two-term model (Occam’s Razor: other things being equal, choose the simplest model.).

So our preferred model is “Sales \sim TV + Radio”.

Case Study 3: Advertising: Viewing the Model



Since this two-term model ignores the contribution of the newspaper channel, the Newspaper spend as a contribution to Sales is just another component of the unmodelled (and apparently random) contribution to Sales.

However, the result is a model where the model “explains” 90% of the variance of the data, which is high for an observational study. **Why? Can we do better?**

Case Study 3: Advertising: Interactions; Interpretation

- Trying powers of the Radio and TV greater than 1 did not offer much more.
- However, by adding the TV, Radio interaction so that the model became “Sales \sim TV + Radio + TV:Radio” or equivalently “Sales \sim TV * Radio”, the loss decreased significantly, indicating the interaction term is valuable.
- All β terms have t –statistic significance of approximately 0.001 which is extremely significant.
- $\beta_0 = 6.75$, $\beta_{\text{TV}} = 0.019$, $\beta_{\text{Radio}} = 0.029$ and $\beta_{\text{TV:Radio}} = 0.001$, indicating that there is a favourable relationship between TV and Radio advertising ($\beta_{\text{TV:Radio}} > 0$), and that additional spending on Radio results in more Sales than the same spending on TV ($\beta_{\text{Radio}} > \beta_{\text{TV}}$).
- Spending on Newspaper advertising should be discontinued as its contribution to Sales is insignificant (indistinguishable from random noise).

Case Study 4: Credit balances - overview

Introducing

- the sklearn approach to regression (we used statsmodels with the Diamonds and Advertising data)
- non-numeric explanatory variables like gender and ethnicity
- more advanced regression modelling, e.g., handling correlated variables

Case Study 4: Credit balances - introduction

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Ethnicity	Balance
1	14.891	3606	283	2	34	11	Male	No	Yes	Caucasian	333
2	106.025	6645	483	3	82	15	Female	Yes	Yes	Asian	903
3	104.593	7075	514	4	71	11	Male	No	No	Asian	580
4	148.924	9504	681	3	36	11	Female	No	No	Asian	964
5	55.882	4897	357	2	68	16	Male	No	Yes	Caucasian	331

Note the presence of some categorical attributes (Gender, Student, Married, Ethnicity). These can participate in linear regression models to predict a numeric response, but must be coded first. For example, Gender can become an indicator (0,1)-valued variable of the form IsFemale. Ethnicity has 3 levels and is replaced by $3-1=2$ indicator variables.





Case Study 4: Credit balances - Removing Data

- the purpose of the analysis is to predict credit balances.
- Basic exploratory data techniques (histograms) soon indicated that there were 2 cohorts
 - ① those who do not use their cards and/or clear their balance each month
 - ② those who use their cards and have nonzero balances
- Removing data relating to the first cohort meant that the remaining data looked more cohesive and also made linear regression easier
- Take-away: look for inconsistent subsets in the data, remove them if possible

Case Study 4: Credit balances - Removing correlated attribute

- Correlations between predictors are relatively high, but that between “Limit” and “Rating” is 1
- Generally, customers with a high rating are allowed to have high credit limits
- Conversely, customers will not be allowed high credit limits unless they have a high credit rating
- “Limit” was removed from the data used for analysis
- Take-away: remove correlated attributes, because they increase the standard error (hence variance) and make the solver’s job much more difficult

Case Study 4: Credit balances - Contribution of Categorical Variables



Which of these categorical attributes has a significant effect on Balance?

Case Study 4: Credit balances - Model building

- Using forward selection as before, the best model was found to be “Balance \sim poly(Income,2) + Rating + Age + Student + Income:Rating”
- Could also use Backward Elimination to prune from a complex model
- For this data, high correlations between features can cause difficulties - we need techniques to handle this

Difficulties caused by correlated features

The Problem : Several features are highly correlated, so the solver has difficulty assigning an importance independently to each.

How it shows up : The condition score is large and several model coefficients take large values with opposite signs. Sometimes the solver gives up.

Solution options :

- ① Remove selected features from the model (simple, does not always work and requires care)
- ② Use *dimensionality reduction* (linear PCA) to derive an uncorrelated subset of the features with least loss in explanatory power (principal components can be opaque)
- ③ Use *regularisation*, to “penalise” large model coefficients (solve a related problem with a different loss function)

Regularisation introduction

Add *regularisation* constraints to make the model work: $\min_{\beta} \|\epsilon\|_2^2 + \lambda p(\beta)$

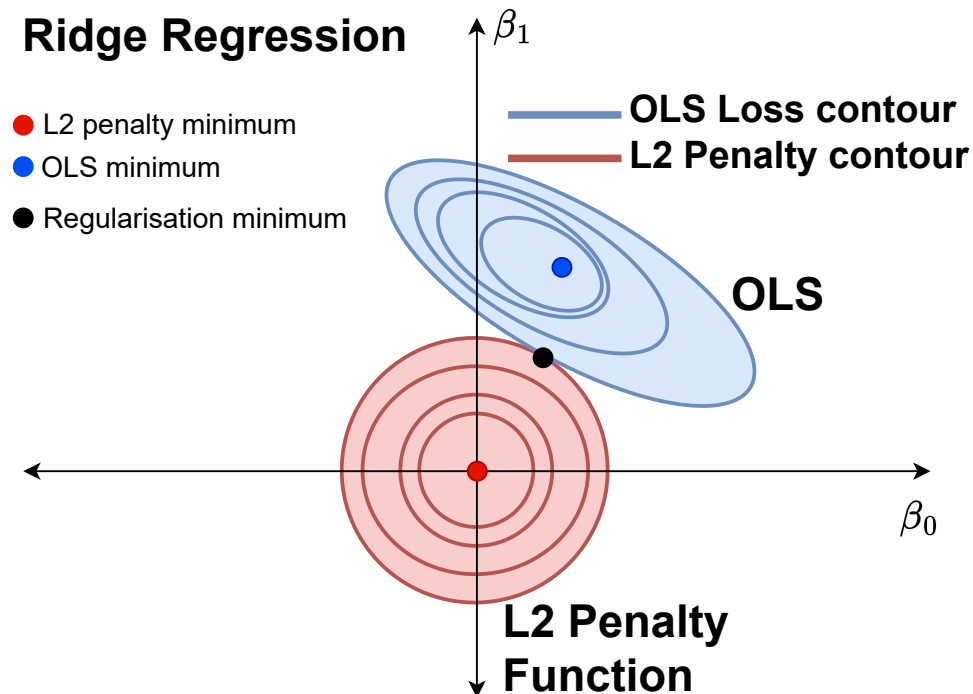
- Options are
 - ① *Ridge Regression* where the penalty term takes the form $p(\beta) = \|\beta\|_2$
 - ② *Lasso* where the penalty term takes the form $p(\beta) = \|\beta\|_1$
- Regularisation has a metaparameter λ - the challenge is to choose a suitable value
 - if too large: tries less to match the data, increases the bias
 - if too small: tries too hard to match the data so $\beta \rightarrow \infty$ and increases the variance

Ridge vs Lasso Regression

Because lasso regression favours the “corners” in parameter space, it tends to set some parameter values to 0 (essentially dropping the associated features). This has the added benefit of making the model smaller and easier to interpret.

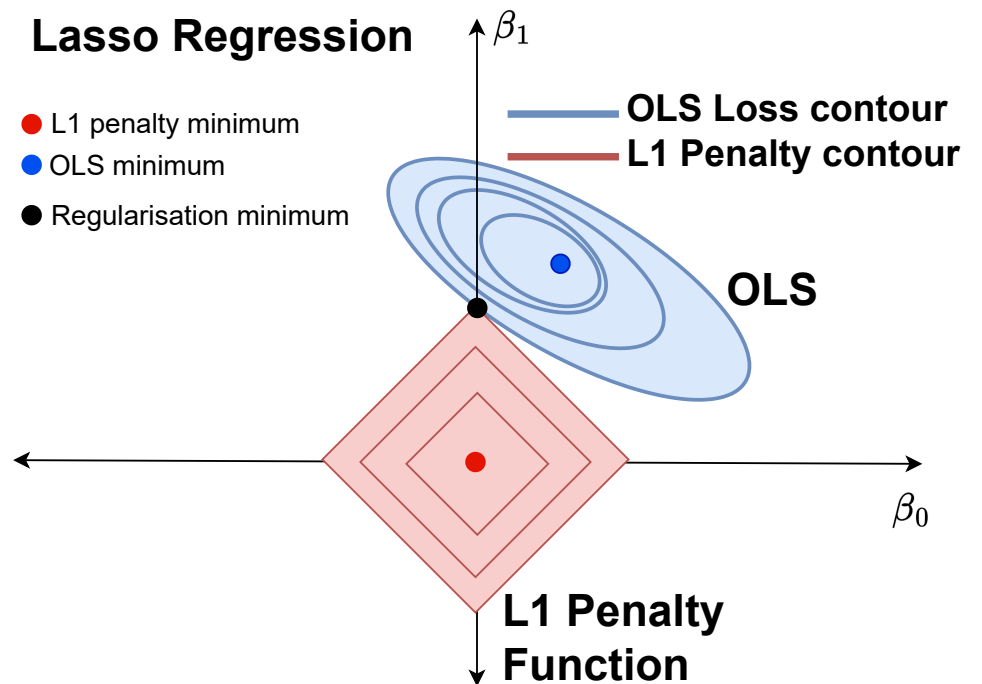
Ridge Regression

Ridge Regression




Lasso Regression


Lasso Regression



Sidebar - vector norms and their unit balls in 2D



By Quartl - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=17428655>



What about unit balls with $0 < p < 1$?

Unit ball, $p = \frac{2}{3}$

Public Domain, <https://commons.wikimedia.org/w/index.php?curid=1770616>

Case Study 4: Credit balances - Regularisation - Searching for λ

- ① Choose a set of candidate λ values
- ② For each candidate λ , use K-fold cross-validation on data subsets to estimate the prediction error for the regularised fit with that λ
- ③ Choose the λ for which the expected error is least
- ④ Now fit all the training data again with this choice of λ

Note that lasso (but not ridge regression) can set particular β_j to 0 (effectively removing them from the model), so it operates more like the *backwards elimination* model building procedure in terms of creating a more frugal model having fewer terms.

Ridge regression downweights certain terms but does not set them to zero. However, it can be more performant, because it keeps some contribution from each feature.

Recall: Attribute independence in Multivariate Data

Definition 3 (Covariance)

$\sigma_{12} = E[(X_1 - \mu_1)(X_2 - \mu_2)]$. In words, for two attributes X_1 and X_2 , with means μ_1 and μ_2 , respectively, σ_{12} is a measure of the linear dependence between them. If they are independent, we can show that $\sigma_{12} = 0$.

Definition 4 ((Variance-)Covariance Matrix)

When there are n numeric attributes, there are $n \times n$ pairs of covariances $\sigma_{ij}, i = 1, \dots, n; j = 1, \dots, n$. The resulting covariance matrix is symmetric and diagonally dominant. This matrix captures the covariance structure of the set of n attributes $\{X_i\}$.


Sometimes it is convenient to work with the correlation matrix, which is a scaled version of the covariance matrix, with elements $\rho_{ij} = \frac{\sigma_{ij}}{\sigma_i \sigma_j}$, which is scaled so that all the diagonal elements are 1 and the off diagonal elements satisfy $-1 < \rho_{ij} < 1$. If two attributes are highly correlated, adding the second into the model does not increase the explanatory power of the model. Therefore, it pays to determine the covariance matrix from the data before building any models.

Feature reduction

- Sometimes it is possible to use intuition to reduce the dimension, by omitting selected attributes.
- Another possibility is to look for groups of correlated attributes (c.f., *mediation*), such as the London and Birmingham measles cases above, and just choose 1 of these.
- More generally, there are techniques that search for a subspace with specified dimension d' of the attributes that captures most of the variance of the full set of attributes having dimension d , where $d' < d$ (often $d' \ll d$).
- The best known of these techniques is *Principal Components Analysis* (PCA).


PCA visualisation

Mapping to 1 Principal Component



Mapping correlated X_1, X_2, X_3 to uncorrelated u_1

Mapping to 2 Principal Components




Mapping correlated X_1, X_2, X_3 to uncorrelated u_1, u_2

PCA interpretation

- Although the data has dimension $d = 3$, it is possible to find the line (on the left; $d = 1$) and plane (on the right, $d = 2$) which retain most of the variance of the data after it has been projected onto this lower dimensional subspace.
- First compute the transformations needed to align the training data with the selected subspace.
- Train the model using the transformed training data and the transformed features (principal components of the original features)
- Can then project other data, e.g., test data, onto the subspace that was derived with the original, training data, and use the model to perform predictions in the transformed space.
- Apply the inverse projection to the data, restoring it to its original orientation. However, because of the use of projections, it is not the same as the original data - the round-trip is “lossy”.
- However, it is helpful to interpret the results in terms of the original attributes.


PCA example




The plot shows that the first 3 **singular values** (associated with principal components u_1, u_2, u_3) capture the bulk of the variance in the training set. Therefore, three attributes, which are transformations of the other 7, are sufficient. You could interpret those attributes as representing the measles outbreaks in three archetypal English cities...

This [youtube video](#) describes PCA concepts well.

Returning to the problematic example




Degree 0 (constant) fit: high bias, low variance



Degree 10 (up to x^{10}) fit: low bias, high variance

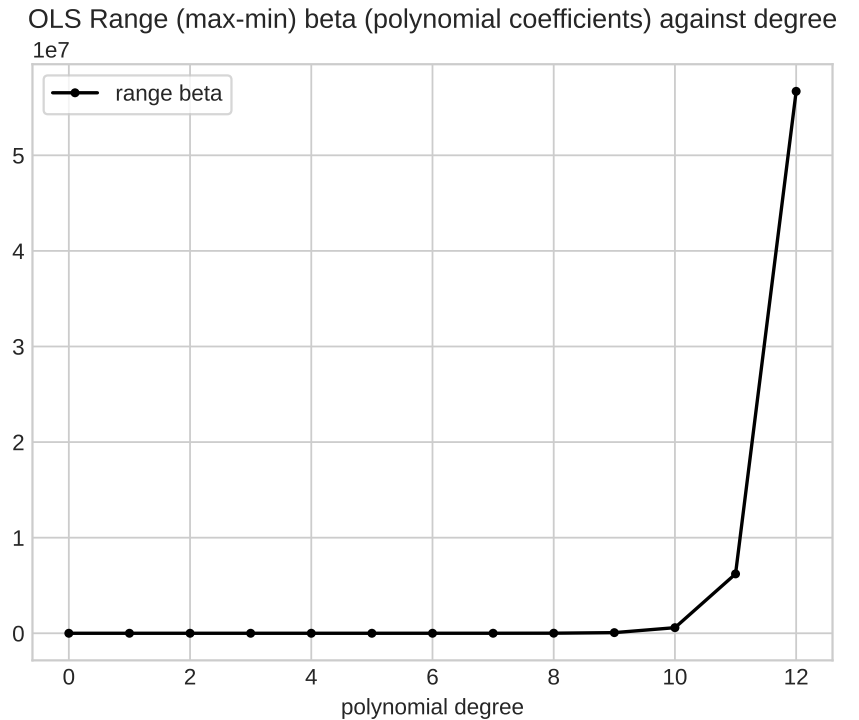
Diagnosis - OLS



Train MSE decreases with degree, Test MSE decreases, then increases

Is there any way we can use high-degree polynomials?


Yes, if we add regularisation...



Polynomial coefficient range (max-min) increases dramatically with degree due to overfitting.


Same data, same features, with regularisation this time

Plot (degree=0) regularisedLR polynomial fit to Actual (30 training, 20 test) points




Degree 0 (constant) fit, $\lambda \approx 0$: no change

Plot (degree=10) regularisedLR polynomial fit to Actual (30 training, 20 test) points




Degree 10 (up to x^{10}) fit: stabilised polynomial

Diagnosis - Regularised Linear Regression



MSE behaviour is affected by choice of λ , but degree 8 or 9 looks good



Polynomial coefficient range (max-min) is controlled - no evidence of overfitting.

So regularisation can control overfitting and/or high correlation between features

Review and summary

- Linear regression is one of the foundations of data mining
- It has two phases, of which the first (learning) is generally the most challenging
- It has many variants, so is quite flexible, but flexibility can be abused!
- Careful validation and model building is essential for success - it is an extension of the exploratory work done earlier in the process
- In machine learning, prediction error is the main focus, but you need to be aware of other considerations such as
 - ① model parsimony (keep them as small as possible!): faster at both training and evaluation time
 - ② the bias-variance dilemma: avoid overfitting and underfitting - remember, your model needs to generalise well from the training to the test set
 - ③ model interpretability: some models are easier to understand because the terms in the model represent concepts from the domain the data is from

Some Additional Resources

- Book: Introduction to Statistical Learning with R (2013) by James, Gareth and Witten, Daniela and Hastie, Trevor and Tibshirani, Robert.
I strongly recommend that you read Chapter 3 of the book, as it is very well written and available online for free.
- Kaggle notebooks relating to the datasets addressed this week. There are many, but searching Kaggle should provide nice examples of data mining in action.
- I wrote a report on linear regression that has been added to moodle.